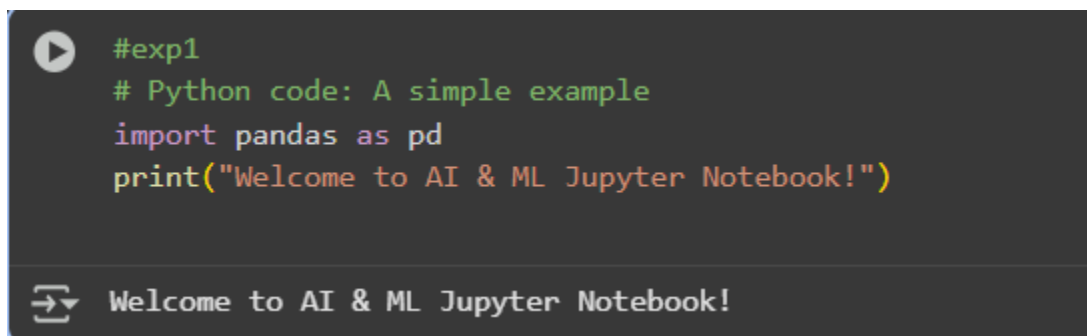


Name: JERON J

Roll no: 231501068

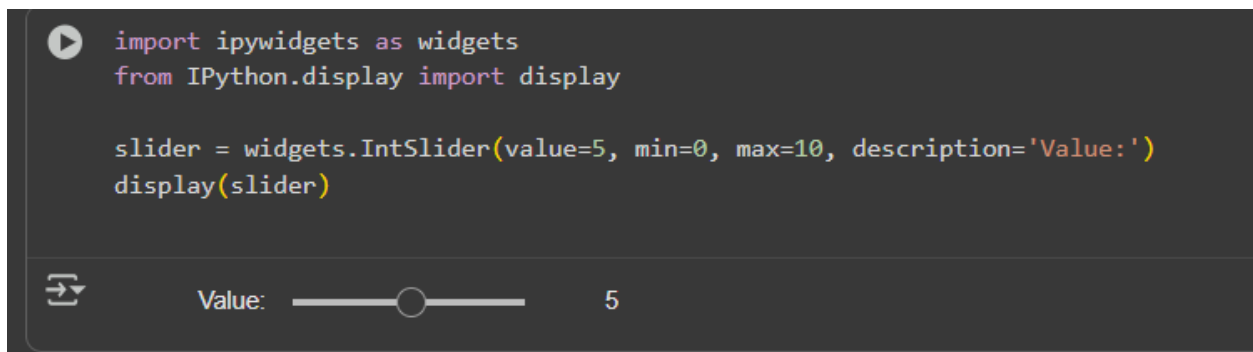
Exp: 1



A screenshot of a Jupyter Notebook cell. The cell contains a play button icon on the left and the following Python code: `#exp1`, `# Python code: A simple example`, `import pandas as pd`, and `print("Welcome to AI & ML Jupyter Notebook!")`. Below the code, there is a separator line and a console output showing the text "Welcome to AI & ML Jupyter Notebook!" with a cursor icon on the left.

```
#exp1
# Python code: A simple example
import pandas as pd
print("Welcome to AI & ML Jupyter Notebook!")
```

Welcome to AI & ML Jupyter Notebook!



A screenshot of a Jupyter Notebook cell. The cell contains a play button icon on the left and the following Python code: `import ipywidgets as widgets`, `from IPython.display import display`, `slider = widgets.IntSlider(value=5, min=0, max=10, description='Value:')`, and `display(slider)`. Below the code, there is a separator line and a console output showing a slider widget. The slider is labeled "Value:" on the left, has a horizontal track with a circular knob in the middle, and the number "5" on the right.

```
import ipywidgets as widgets
from IPython.display import display

slider = widgets.IntSlider(value=5, min=0, max=10, description='Value:')
display(slider)
```

Value:

Exp: 2

Code:

```
import pandas as pd
from sqlalchemy import create_engine
import sqlite3

# 1. Import CSV from URL
csv_url = "https://raw.githubusercontent.com/datasciencedojo/datasets/master/titanic.csv"
df_csv = pd.read_csv(csv_url)
print("CSV Data (Titanic) sample:")
print(df_csv.head())

# 2. Create and save a small Excel file locally for demo (you can replace this with your own Excel file)
df_csv.head(10).to_excel("sample_titanic.xlsx", index=False)

# 3. Import Excel file you just created
df_excel = pd.read_excel("sample_titanic.xlsx")
print("\nExcel Data sample:")
print(df_excel.head())

# 4. Create a simple SQLite DB in memory and load CSV data into it for demo
conn = sqlite3.connect(':memory:')
df_csv.to_sql('titanic', conn, index=False, if_exists='replace')

# Query SQL table
query = "SELECT * FROM titanic LIMIT 5"
df_sql = pd.read_sql(query, conn)
print("\nSQL Data sample:")
print(df_sql)

# 5. Web scraping: Get first table from a Wikipedia page
url = 'https://en.wikipedia.org/wiki/List_of_countries_by_GDP_(nominal)'
tables = pd.read_html(url)
print(f"\nNumber of tables scraped from webpage: {len(tables)}")

df_web = tables[0]
print("\nWeb scraped Data sample:")
print(df_web.head())

# 6. Export the Titanic CSV DataFrame to Excel
df_csv.to_excel('exported_titanic.xlsx', index=False)
print("\nExported Titanic DataFrame to 'exported_titanic.xlsx'")
```

Output:

```

0      Parch      Ticket      Fare Cabin Embarked
1      0      A/5 21171      7.2500 NaN      S
2      0      PC 17599      71.2833 C85      C
3      0 STON/O2. 3101282      7.9250 NaN      S
4      0      113803      53.1000 C123      S
5      0      373450      8.0500 NaN      S

Excel Data sample:
  PassengerId  Survived  Pclass  \
0             1         0       3
1             2         1       1
2             3         1       3
3             4         1       1
4             5         0       3

                                Name      Sex  Age  SibSp  \
0                                Braund, Mr. Owen Harris    male  22.0      1
1  Cumings, Mrs. John Bradley (Florence Briggs Th...  female  38.0      1
2                                Heikkinen, Miss. Laina  female  26.0      0
3      Futrelle, Mrs. Jacques Heath (Lily May Peel)  female  35.0      1
4                                Allen, Mr. William Henry    male  35.0      0

  Parch      Ticket      Fare Cabin Embarked
0      0      A/5 21171      7.2500 NaN      S
1      0      PC 17599      71.2833 C85      C
2      0 STON/O2. 3101282      7.9250 NaN      S
3      0      113803      53.1000 C123      S
4      0      373450      8.0500 NaN      S

SQL Data sample:
  PassengerId  Survived  Pclass  \
0             1         0       3
1             2         1       1
2             3         1       3
3             4         1       1
4             5         0       3

                                Name      Sex  Age  SibSp  \
0                                Braund, Mr. Owen Harris    male  22.0      1
1  Cumings, Mrs. John Bradley (Florence Briggs Th...  female  38.0      1
2                                Heikkinen, Miss. Laina  female  26.0      0
3      Futrelle, Mrs. Jacques Heath (Lily May Peel)  female  35.0      1
4                                Allen, Mr. William Henry    male  35.0      0
```

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	None	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	None	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	None	S

Number of tables scraped from webpage: 7

Web scraped Data sample:

0 Largest economies in the world by GDP (nominal...

Exported Titanic DataFrame to 'exported_titanic.xlsx'

Exp:3

Code:

```
import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler, MinMaxScaler

# Load the dataset
file_path = '/content/logistic_regression_dataset.csv'
df = pd.read_csv(file_path)

# 1. Handling missing values
print("Missing values per column:\n", df.isnull().sum())
for col in df.columns:
    if df[col].dtype in ['int64', 'float64']:
        df[col].fillna(df[col].mean(), inplace=True)
    else:
        df[col].fillna(df[col].mode()[0], inplace=True)

df.dropna(thresh=len(df.columns) - 1, inplace=True) # Drop rows with too many missing

# 2. Remove duplicates and unnecessary columns
df.drop_duplicates(inplace=True)
df = df.loc[:, ~df.columns.str.contains('^Unnamed')] # Remove unnamed columns
if 'ID' in df.columns:
    df.drop(columns=['ID'], inplace=True)

# 3. Data type conversion
for col in df.select_dtypes(include='object').columns:
    if df[col].nunique() < 50:
        df[col] = df[col].astype('category')

for col in df.columns:
    if df[col].dtype == 'object':
        try:
            df[col] = pd.to_numeric(df[col])
        except:
            pass
```

```

# 4. Normalize numeric data
numeric_cols = df.select_dtypes(include=['int64', 'float64']).columns

scaler_std = StandardScaler()
df_standardized = df.copy()
df_standardized[numeric_cols] = scaler_std.fit_transform(df[numeric_cols])

scaler_mm = MinMaxScaler()
df_minmax = df.copy()
df_minmax[numeric_cols] = scaler_mm.fit_transform(df[numeric_cols])

# Show first 5 rows of cleaned data
print("Cleaned Data Sample:")
print(df.head())

print("\nStandardized Data Sample:")
print(df_standardized.head())

print("\nMin-Max Scaled Data Sample:")
print(df_minmax.head())

```

Output:

```

Missing values per column:
[ ] Age      0
    Income   0
    Credit_Score  0
    Loan_Amount  0
    Employment_Years  0
    Loan_Status  0
    dtype: int64
Cleaned Data Sample:
   Age  Income  Credit_Score  Loan_Amount  Employment_Years  Loan_Status
0   60     11         710           39             20             0
1   27      5         574           38              1             0
2   21     13         367            5              7             0
3   37      3         516           43              2             0
4   35      2         622           25              1             0

Standardized Data Sample:
   Age      Income  Credit_Score  Loan_Amount  Employment_Years  \
0  1.708662  0.549104    0.781864    0.910189             1.701140
1 -0.948769 -0.889026   -0.086229    0.836463            -1.379781
2 -1.431938  1.028481   -1.407518   -1.596518            -0.406858
3 -0.143487 -1.368403   -0.456445    1.205096            -1.217627
4 -0.304543 -1.608091    0.220157   -0.121984            -1.379781

```

```

Loan_Status
0    -0.442326
1    -0.442326
2    -0.442326
3    -0.442326
4    -0.442326

Min-Max Scaled Data Sample:
   Age   Income  Credit_Score  Loan_Amount  Employment_Years  Loan_Status
0  1.000  0.692308    0.745421    0.755556           1.00         0.0
1  0.175  0.230769    0.496337    0.733333           0.05         0.0
2  0.025  0.846154    0.117216    0.000000           0.35         0.0
3  0.425  0.076923    0.390110    0.844444           0.10         0.0
4  0.375  0.000000    0.584249    0.444444           0.05         0.0

/tmp/ipython-input-6-2547170765.py:16: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

df[col].fillna(df[col].mean(), inplace=True)

```

Exp: 4

Code:

```

import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

# Step 2: Load Wine Quality Red dataset from UCI (semicolon-separated)
url = 'https://archive.ics.uci.edu/ml/machine-learning-databases/wine-quality/winequality-red.csv'
df = pd.read_csv(url, sep=';')

# Step 3: Quick look at the data
print("First 5 rows:")
print(df.head())

print("\nShape of dataset:", df.shape)

print("\nColumn names:")
print(df.columns)

print("\nInfo about data types and null values:")
print(df.info())

print("\nSummary statistics:")
print(df.describe())

# Step 4: Filtering and subsetting examples
print("\nRows where quality >= 7:")
print(df[df['quality'] >= 7].head())

# Select subset columns
print("\nSelect columns: alcohol and quality")
print(df[['alcohol', 'quality']].head())

```

```
# Step 5: Descriptive statistics for 'alcohol' and 'quality'
print("\nDescriptive statistics for 'alcohol':")
print(f"Mean: {df['alcohol'].mean():.2f}")
print(f"Median: {df['alcohol'].median():.2f}")
print(f"Mode: {df['alcohol'].mode().values}")

print("\nDescriptive statistics for 'quality':")
print(f"Range: {df['quality'].max() - df['quality'].min()}")
print(f"Variance: {df['quality'].var():.2f}")
print(f"Standard Deviation: {df['quality'].std():.2f}")

# Step 6: Visualizations

# Histogram for alcohol content
plt.figure(figsize=(8,4))
sns.histplot(df['alcohol'], kde=True, bins=30)
plt.title('Alcohol Content Distribution')
plt.xlabel('Alcohol')
plt.ylabel('Frequency')
plt.show()

# Boxplot of quality scores
plt.figure(figsize=(6,4))
sns.boxplot(x='quality', data=df)
plt.title('Boxplot of Wine Quality')
plt.xlabel('Quality Score')
plt.show()

# Correlation heatmap
plt.figure(figsize=(12,8))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Heatmap of Wine Quality Dataset')
plt.show()
```

Output:

```
First 5 rows:
  fixed acidity  volatile acidity  citric acid  residual sugar  chlorides \
0          7.4           0.70         0.00         1.9         0.076
1          7.8           0.88         0.00         2.6         0.098
2          7.8           0.76         0.04         2.3         0.092
3         11.2           0.28         0.56         1.9         0.075
4          7.4           0.70         0.00         1.9         0.076

  free sulfur dioxide  total sulfur dioxide  density  pH  sulphates \
0             11.0             34.0  0.9978  3.51         0.56
1             25.0             67.0  0.9968  3.20         0.68
2             15.0             54.0  0.9970  3.26         0.65
3             17.0             60.0  0.9980  3.16         0.58
4             11.0             34.0  0.9978  3.51         0.56

  alcohol  quality
0      9.4        5
1      9.8        5
2      9.8        5
3      9.8        6
4      9.4        5

Shape of dataset: (1599, 12)

Column names:
Index(['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar',
      'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'density',
      'pH', 'sulphates', 'alcohol', 'quality'],
      dtype='object')

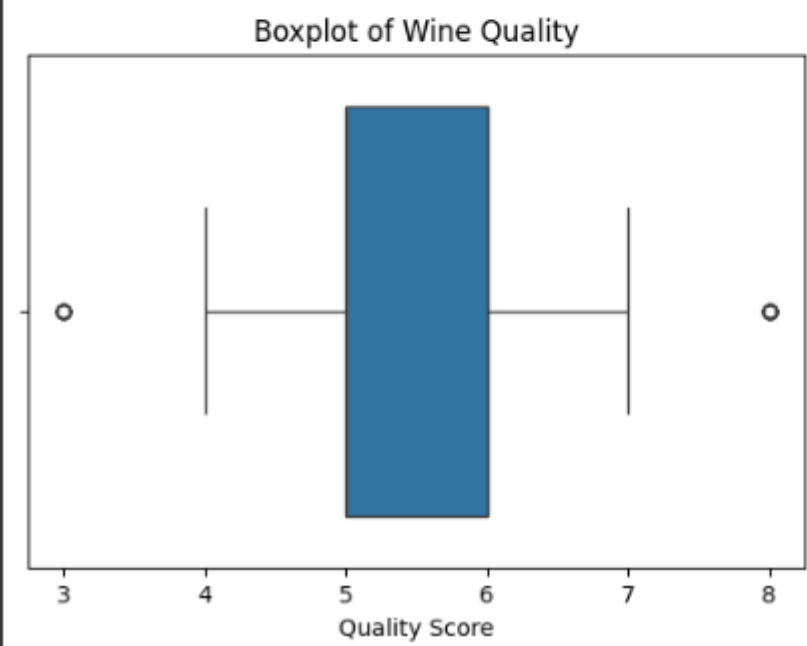

Info about data types and null values:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1599 entries, 0 to 1598
```


```
Select columns: alcohol and quality
  alcohol  quality
0      9.4        5
1      9.8        5
2      9.8        5
3      9.8        6
4      9.4        5

Descriptive statistics for 'alcohol':
Mean: 10.42
Median: 10.20
Mode: [9.5]

Descriptive statistics for 'quality':
Range: 5
Variance: 0.65
Standard Deviation: 0.81
```


Range: 5
Variance: 0.65
Standard Deviation: 0.81





Exp: 5

Code:

```
import pandas as pd
import matplotlib.pyplot as plt

# Load the dataset from the GitHub Gist
url = 'https://gist.github.com/svlsm1/3b5dd6723510ae43a0f121213f6583fc/raw/StudentsPerformance.csv'
df = pd.read_csv(url)

# Preview data
print(df.head())

# Prepare data
x = df.index # Use row index for line plot
y_math = df['math_score']
y_reading = df['reading_score']
```

```

# For bar chart: average scores by gender
avg_by_gender = df.groupby('gender')[['math_score', 'reading_score', 'writing_score']].mean()

# For histogram: distribution of math scores
hist_data = df['math_score']

# Generate subplots
fig, axs = plt.subplots(1, 3, figsize=(18, 5))

# 1. Line chart: Math vs Reading
axs[0].plot(x, y_math, label='Math Score', alpha=0.7)
axs[0].plot(x, y_reading, label='Reading Score', alpha=0.7)
axs[0].set_title('Line Chart: Math vs Reading Scores')
axs[0].set_xlabel('Sample Index')
axs[0].set_ylabel('Score')
axs[0].legend()
axs[0].grid(True)

# 2. Bar chart: Average by gender
avg_by_gender.plot(kind='bar', ax=axs[1])
axs[1].set_title('Average Scores by Gender')
axs[1].set_xlabel('Gender')
axs[1].set_ylabel('Average Score')
axs[1].legend(title='Subject')
axs[1].grid(axis='y')

# 3. Histogram: Math score distribution
axs[2].hist(hist_data, bins=20, color='purple', edgecolor='black')
axs[2].set_title('Histogram: Math Score Distribution')
axs[2].set_xlabel('Math Score')
axs[2].set_ylabel('Frequency')

# Adjust layout and show
plt.tight_layout()
plt.show()

```


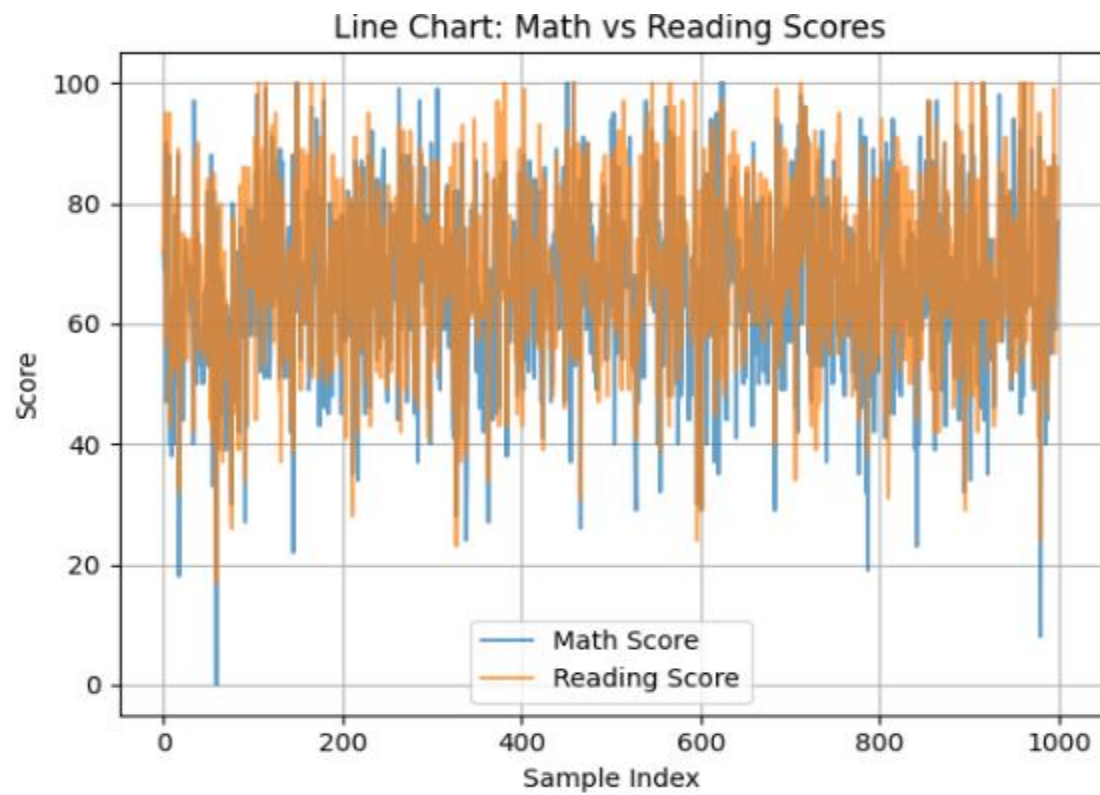
Output:

```

➡ gender race_ethnicity parental_level_of_education lunch \
0 female group B bachelor's degree standard
1 female group C some college standard
2 female group B master's degree standard
3 male group A associate's degree free/reduced
4 male group C some college standard

test_preparation_course math_score reading_score writing_score
0 none 72 72 74
1 completed 69 90 88
2 none 90 95 93
3 none 47 57 44
4 none 76 78 75

```



Histogram: Math Score Distribution

